

# CSC 411: Lecture 3 - Linear Classification

Ethan Fetaya, James Lucas and Emad Andrews

This lecture:

- Linear classification (binary).
- First order optimization.
- Key concepts:
  - Decision boundaries.
  - Loss functions.
  - metrics to evaluate classification.
  - Stochastic gradient descent.

Last week: Mapping  $\mathbf{x} \in \mathbb{R}^d$  into  $y \in \mathbb{R}$ .

This week: Mapping  $\mathbf{x} \in \mathbb{R}^d$  into categorical  $y$  (in a finite set  $S$ ).

Usually use  $S = \{1, \dots, k\}$ ,  $S = \{0, 1\}$  or  $S = \{-1, 1\}$  (our focus now).

Linear model:  $\hat{y} = f(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \mathbf{x}$  outputs a real score.

How do we turn it into a binary decision? Threshold -

$$\hat{y} = f(\mathbf{x}, \mathbf{w}) = \text{sign}(\mathbf{w}^T \mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} \geq 0 \\ -1 & \text{if } \mathbf{w}^T \mathbf{x} < 0 \end{cases}$$

Decision boundary is the hyperspace defined by  $\mathbf{w}$ .

$\mathbf{w}^T \mathbf{x} = 0$  is a hyperplane (line in  $d = 2$ ) passing through the origin and orthogonal to  $\mathbf{w}$ .  $\mathbf{w}^T \mathbf{x} + w_0 = 0$  shifts it by  $w_0$ .




Figure from G. Shakhnarovich




Decision boundary is invariant to scaling.

If we can separate the classes by a hyperplane, the problem is **linearly separable**




Causes of non perfect separation:

- Model is too simple.
- Noise (optimal classifier might not be perfect).
- Errors in data targets (miss labelings).
- Simple features that do not account for all variations.
- Need different feature parametrization.



X,Y coordinates




Polar coordinates

Should we make the model complex enough to have perfect separation in the training data?

Learning consists of finding a good **decision boundary**.

We need to find  $\mathbf{w}$  (direction) and  $w_0$  (location) of the boundary.

What does "good" mean? Is this boundary good?



We need a criteria that tell us how to select the parameters.

A natural loss function: **zero-one loss**.  $\ell_{0-1}(\hat{y}, y) = \begin{cases} 1 & \text{if } y \neq \hat{y} \\ 0 & \text{if } y = \hat{y} \end{cases}$

Is this minimization easy to do? Why?

**Asymmetric Binary Loss:** Should we treat both types of mistakes

equally?  $\ell_{ABL}(\hat{y}, y) = \begin{cases} \alpha & \text{if } y = 0 \wedge \hat{y} = 1 \\ \beta & \text{if } y = 1 \wedge \hat{y} = 0 \\ 0 & \text{if } y = \hat{y} \end{cases}$

When is this important?

Goal: Optimizing  $\ell_{0-1}$  (or  $\ell_{ABL}$ ).

Problem: (NP)hard, piecewise constant.

Approach: use a **surrogate loss**  $\hat{\ell}$  to optimize instead.


What makes a good surrogate loss?

- Easy to optimize
  - (Piecewise) Smooth.
  - Convex.
- Representative - low surrogate loss means low original loss.
  - Upper bound  $\forall y \forall \hat{y} \ell(y, \hat{y}) \leq \hat{\ell}(y, \hat{y})$ .



Is  $\ell_2(y, \hat{y}) = (y - \hat{y})^2$  loss a good surrogate?

Easy to optimize? ✓ Representative? so-so



We will see better surrogates soon.

How to evaluate how good my classifier is? [Metrics](#)

- Metrics on a dataset is what we care about (performance).
- We typically cannot directly optimize for the metrics.
- Our loss function should reflect the problem we are solving. We then hope it will yield models that will do well on our dataset.

**Accuracy:** Percent of correct predictions,  $1 - \ell_{0-1}(w)$ .

Is it a good measure? Data balanced? Unbalanced?

**Recall:** The fraction of relevant instances that are retrieved.

$$R = \frac{TP}{TP + FN} = \frac{TP}{\text{all groundtruth instances}}$$


**Precision:** The fraction of retrieved instances that are correct.

$$P = \frac{TP}{TP + FP} = \frac{TP}{\text{all positive predictions}}$$

**F1 score:** Harmonic mean of precision and recall.

$$F1 = 2 \frac{P \cdot R}{P + R}$$


**Precision-Recall curve:** Trade-off between recall and precision using the decision threshold.




**Average Precision (AP):** area under the curve.

We might be interest in a single working point (recall or precision).

Receiver Operator Characteristic (ROC): Trade-off between false-positive-rate (**FPR**) and true-positive-rate (**TPR**) using the decision threshold.



(a) Comparison in ROC space



(b) Comparison in PR space

Better in ROC  $\Rightarrow$  better in PR (not always vice-versa).

Difference can be big with unbalanced data

---

<sup>1</sup>Figure from "The Relationship Between Precision-Recall and ROC Curves"


Once we decide on a (smooth) loss  $\ell$  - how do we find

$$\mathbf{w} = \arg \min L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ell(y_i, f(x_i, \mathbf{w}))?$$

One straightforward method: gradient descent

- initialize  $\mathbf{w}_0$  (e.g., randomly)
- repeatedly update  $\mathbf{w}$  based on the gradient

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \lambda_t \nabla_{\mathbf{w}} L(\mathbf{w}_t)$$




$\lambda$  is the learning rate.

Update rule:  $\mathbf{w}_{t+1} = \mathbf{w}_t - \lambda_t \nabla_{\mathbf{w}} L(\mathbf{w}_t)$

Finding a good learning rate is very important.

- Too large  $\lambda$ : unstable and can diverge.
- Too low  $\lambda$ : stable but very slow progress.
- Line search methods - usually too slow.
- Standard to decay  $\lambda$  as learning progresses.



Commonly found using simple grid search, some automatic tools exist.

---

<sup>1</sup>Image credit: <https://www.slideshare.net/simaokasonse/learning-deep-learning>.

What is the computational cost of computing  $\nabla_{\mathbf{w}} L(\mathbf{w}_t)$ ?

$L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ell(y_i, f(x_i, \mathbf{w}))$  - grows linearly in  $N$  (number of data points).

Huge (millions/billions) dataset  $\Rightarrow$  large cost for a tiny update!

Solution: Stochastic gradient descent. Instead of computing gradient  $g_t = \nabla_{\mathbf{w}} L(\mathbf{w}_t) = \frac{1}{N} \sum_{i=1}^N \nabla \ell(y_i, f(x_i, \mathbf{w}))$ , pick random datum  $j$  and compute  $\hat{g}_t = \nabla \ell(y_j, f(x_j, \mathbf{w}))$

Will it work? Theoretically - yes (with the right learning rate decay).  
Practically - very noisy.



Better solution: Mini-batch. Middle-ground, average  $1 < m \ll N$  gradients.

Mean is still  $g_t$  but variance is lower. Trade-off between accuracy (big batch) and runtime (small batch).

---

**Algorithm 1** Mini-batch gradient descent epoch

---

- 1: Randomly shuffle examples in the training set
- 2: **for**  $i = 0$  to  $N/m$  **do**
- 3:     Update:

$$\mathbf{w} \leftarrow \mathbf{w} + \frac{1}{m} \sum_{j=0}^{m-1} \nabla \ell(y^{m \cdot i + j}, f(x^{m \cdot i + j}, \mathbf{w}))$$

- 4: **end for**
- 

This simple idea is an important component behind a lot of recent success.

People commonly use the term SGD for mini-batch optimization.