

How to Elect a Leader Faster than a Tournament

Dan Alistarh, Rati Gelashvili, Adrian Vladu
PODC 2015

Presenter: Lily Li

CSC 2221: Introduction to Distributed Algorithms


30 November 2017

Results New and Old

Tournament Tree

Step Complexity: $O(\log n)$

Message Complexity: $O(n^2 \log n)$



Results New and Old

Tournament Tree


Step Complexity: $O(\log n)$

Message Complexity: $O(n^2 \log n)$

Poison Pill

Step Complexity: $O(\log^* n)$

Message Complexity: $O(n^2)$



Results New and Old

Tournament Tree

Step Complexity: $O(\log n)$


Message Complexity: $O(n^2 \log n)$

Poison Pill


Step Complexity: $O(\log^* n)$

Message Complexity: $O(n^2)$


Lower bound: $\Omega(n^2)$




Algorithm Overview




Algorithm Overview




Algorithm Overview




Algorithm Overview




Algorithm Overview




Algorithm Overview



Algorithm Overview



Algorithm Overview



Local Variables and Communication Primitives

Local variables of process p_j :

- S_i : records the states of all observed processes
- V_i : view stores the state vectors received from other processes

Local Variables and Communication Primitives

Local variables of process p_j :

- S_i : records the states of all observed processes
- V_i : view stores the state vectors received from other processes

High-Level Protocol **COMMUNICATE**< p > (where p is a procedure which broadcasts a message and wait for a response):

Executes p and waits for $\left\lceil \frac{n}{2} \right\rceil + 1$ processes to respond before proceeding.

p can be `broadcast(v)` or `collect()` where

1. If p_j receives `broadcast(v)` from p_i , store $S_j[i] \leftarrow v$ and send $p_i \langle ack \rangle$.
2. If p_j receives `collect()` from p_i , send $p_i \langle S_j \rangle$.

Local Variables and Communication Primitives

Local variables of process p_j :

- S_i : records the states of all observed processes
- V_i : view stores the state vectors received from other processes

High-Level Protocol **COMMUNICATE** $\langle p \rangle$ (where p is a procedure which broadcasts a message and wait for a response):

Executes p and waits for $\left\lceil \frac{n}{2} \right\rceil + 1$ processes to respond before proceeding.

p can be `broadcast(v)` or `collect()` where

1. If p_j receives `broadcast(v)` from p_i , store $S_j[i] \leftarrow v$ and send $p_i \langle ack \rangle$.
2. If p_j receives `collect()` from p_i , send $p_i \langle S_j \rangle$.

`random(r)` : outputs 1 with probability r and 0 with probability $1-r$.

(Simple) Poison Pill

Algorithm 1. Code for process p_i during one phase.

```
1: Initialize:  $S_i = [w, \dots, w]$ ,  $V_i = \text{EMPTY}$ 
2:    $S_i[i] = c$ 
3:   COMMUNICATE<broadcast( $S_i[i]$ )>
4:    $S_i[i] = \text{random}(1/\text{sqrt}(n))$ 
5:   COMMUNICATE<broadcast( $S_i[i]$ )>
6:    $V_i = \text{COMMUNICATE}$ <collect()>
7:   if( $S_i[i] = 0$  and exists  $k$ :  $p_k$  is active)
8:     return 0
9:   return 1
```


(Simple) Poison Pill

Algorithm 1. Code for process p_i during one phase.

```
1: Initialize:  $S_i = [w, \dots, w]$ ,  $V_i = \text{EMPTY}$ 
2:    $S_i[i] = c$ 
3:   COMMUNICATE<broadcast( $S_i[i]$ )>
4:    $S_i[i] = \text{random}(1/\text{sqrt}(n))$ 
5:   COMMUNICATE<broadcast( $S_i[i]$ )>
6:    $V_i = \text{COMMUNICATE}$ <collect()>
7:   if ( $S_i[i] = 0$  and exists  $k$ :  $p_k$  is active)
8:     return 0
9:   return 1
```

Claim 1. If all processes return then at least one will output 1.

(Simple) Poison Pill

Algorithm 1. Code for process p_i during one phase.

```
1: Initialize:  $S_i = [w, \dots, w]$ ,  $V_i = \text{EMPTY}$ 
2:    $S_i[i] = c$ 
3:   COMMUNICATE<broadcast( $S_i[i]$ )>
4:    $S_i[i] = \text{random}(1/\text{sqrt}(n))$ 
5:   COMMUNICATE<broadcast( $S_i[i]$ )>
6:    $V_i = \text{COMMUNICATE}$ <collect()>
7:   if ( $S_i[i] = 0$  and exists  $k$ :  $p_k$  is active)
8:     return 0
9:   return 1
```

Claim 1. If all processes return then at least one will output 1.

Claim 2. The expected number of processes which output 1 is $O(\sqrt{n})$.

(Simple) Poison Pill

Algorithm 1. Code for process p_i during one phase.

```
1: Initialize:  $S_i = [w, \dots, w]$ ,  $V_i = \text{EMPTY}$ 
2:    $S_i[i] = c$ 
3:   COMMUNICATE<broadcast( $S_i[i]$ )>
4:    $S_i[i] = \text{random}(1/\text{sqrt}(n))$ 
5:   COMMUNICATE<broadcast( $S_i[i]$ )>
6:    $V_i = \text{COMMUNICATE}$ <collect()>
7:   if ( $S_i[i] = 0$  and exists  $k$ :  $p_k$  is active)
8:     return 0
9:   return 1
```

Claim 1. If all processes return then at least one will output 1.

Claim 2. The expected number of processes which output 1 is $O(\sqrt{n})$.

Lemma 1. If a process receives 1 from `random` at time t , then all processes which received 0 from `random` at any time $\geq t$ will return 0.

...to Recap

- processes commit to an action before performing the action
- the probability of surviving each round changes dynamically

...to Recap

- processes commit to an action before performing the action
- the probability of surviving each round changes dynamically

...the Result

$O(\log^* n)$ step complexity and
 $O(n^2)$ message complexity.

...to Recap

- processes commit to an action before performing the action
- the probability of surviving each round changes dynamically

...the Result

$O(\log^* n)$ step complexity and
 $O(n^2)$ message complexity.

This randomized algorithm can be used to solve the tight renaming problem in $O(\log^2 n)$ steps using $O(n^2)$ messages.