

# Alistarh et al. (2015) — “How to Elect a Leader Faster than a Tournament”

November 21, 2017; rev. November 23, 2017

Lily Li

## 1 Introduction

**Problem Definition.** An algorithm which solves leader election among  $n$  processes must satisfy: (1) termination - every process  $p_i$  must eventually output *win* or *lose* and (2) unique winner - there is exactly one process which outputs *win*. Further, no process may lose before the eventual winner starts its execution and the algorithm must be linearizable; the first operation is *win* and all subsequent operations is *lose*. In-class, we have seen an leader election algorithm which builds a *tournament tree* with  $n$  leaves. The step and message complexities of this decades old algorithm are  $\Theta(\log n)$  and  $\Theta(n^2 \log n)$  respectively.

**Model.** We have a complete asynchronous message passing system with at most  $f < \lceil n/2 \rceil$  faulty processes. Further we assume a strong adversary which can examine the systems state and the outcome of random coin flips adjusting the schedules of steps, message delivery events, and failures.

**Results.** In this work by Alistarh et al. presented at PODC 2015, the authors present a leader election algorithm which requires  $O(\log^* n)$  steps by each process and  $O(n^2)$  point-to-point messages. They pair this with a lower bound of  $\Omega(n^2)$  message, proving optimal message complexity.

## 2 Leader Election Algorithm

In the code we will make use of the high-level protocol COMMUNICATE. Let  $p$  be a procedure which consists of broadcasting a message and waiting for a response. COMMUNICATE( $p$ ) will wait for at least  $\lceil n/2 \rceil + 1$  responses before proceeding (this is reminiscent to what we did when we tried to simulate single-reader single-writer registers in a message passing system). We call a set of  $\lceil n/2 \rceil + 1$  processes a **quorum** since these arise often in our analysis. A process can execute the following operations:

**broadcast( $v$ ):** process  $p_i$  broadcasts value  $v$  to all other processes (including itself). If process  $p_j$  receives broadcast  $v$  from  $p_i$ , it updates  $S_j[i] \leftarrow v$ .

**collect():** process  $p_i$  sends a collect message to all other processes (including itself). If process  $p_j$  receives collect() from  $p_i$ , it sends message  $\langle S_j \rangle$  to  $p_i$ .

**random( $p$ ):** (local operation) process  $p_i$  receives 1 with probability  $p$  and 0 otherwise.

What we want is an algorithm which iteratively reduces the number of participating processes. At each round processes flip a random coin and output *lose* or *survive* depending on the result. Surviving processes will try again in the next round. A naïve implementation would fail since the adversary can schedule the “losing” processes first.

### 2.1 Poison Pill (Homogeneous)

See Algorithm 1. Process  $p_i$  has two local variables: a vector  $S_i$  of length  $n$  which record the states of processes (with all entries initially set to  $w$ ) and an  $n \times n$  matrix  $V_i$  storing the view as seen by other processes (initially empty). States can be any of  $w$  (waiting),  $c$  (committed), 0 (lose), and 1 (survive). A process  $p_j$  is **active** from the perspective of  $p_i$  if column  $j$  in  $V_i$  contains at least one entry of status 1 or  $c$  and no entries of status 0.

**Claim 1.** *If all processes return then some process outputs 1.*

*Proof.* Suppose for a contradiction that all processes output 0. If a process receives 1 upon executing  $\text{random}(\frac{1}{\sqrt{n}})$ , then it must output 1 so all processes received 0. Let process  $p_i$  be the last to complete the broadcast operation on Line 5. When  $p_i$ ’s broadcast operation finishes, the 0 of every process gets stored by a quorum. Consider  $V_i$  at the completion of Line 6. Every column must have one 0-entry since the

**Algorithm 1** Homogeneous Poison Pill: code for process  $p_i$ .

---

```

1: Initialize  $S_i[1..n] = [w, \dots, w]$  and  $V_i[1..n][1..n] = \emptyset$ 
2:  $S_i[i] \leftarrow c$ 
3: COMMUNICATE( $\text{broadcast}(S_i[i])$ )
4:  $S_i[i] \leftarrow \text{random}\left(\frac{1}{\sqrt{n}}\right)$ 
5: COMMUNICATE( $\text{broadcast}(S_i[i])$ )
6:  $V_i \leftarrow \text{COMMUNICATE}(\text{collect}())$ 
7: if  $S_i[i] = 0$  and  $\exists k : p_k$  is active then
8:   return 0
9: end if
10: return 1

```

---

set of processes that stored 0 must overlap with the set of processes which successfully sent their status vector to  $p_i$ . Thus  $p_i$  must output 1.  $\square$

**Claim 2.** The expected number of processes that return 1 is  $O(\sqrt{n})$ .

*Proof.* Suppose  $p_i$  received 1 from its execution of random at  $t_i$ . Then all process which received 0 at any time  $\geq t_i$  must output 0. Consider such a process  $p_j$ . Observe that at  $t_i$ ,  $S_i[i] = \top$  has been broadcast to a quorum. Thus the set of processes which received  $S_i[i]$  and the set of processes which sends their status to  $p_j$  upon its execution of collect() overlaps.  $p_j$  will see that  $p_i$  has  $\top$  or 1 and will return 0. Now simply observe that the expected number of 1s as well as the expected index for the first 1 are both  $\sqrt{n}$ .  $\square$

## 2.2 Poison Pill (Heterogeneous)

It can be shown that there will be  $\Omega(\sqrt{n})$  survivors at every round of the homogeneous algorithm — the probability of getting 1 is fixed. To improve this bound we need to change this probability.

**Algorithm 2** Heterogeneous Poison Pill: code for process  $p_i$ .

---

```

1:  $S_i[i] \leftarrow \{c, \{\}\}$ 
2: COMMUNICATE( $\text{broadcast}(S_i[i])$ )
3:  $V_i \leftarrow \text{COMMUNICATE}(\text{collect}())$ 
4:  $t \leftarrow \{j : \exists k : V_i[k][j] \neq \perp\}$ 
5:  $p \leftarrow \frac{\log |t|}{|t|}$  #  $p \leftarrow 1$  if  $|t| = 1$ 
6:  $S_i[i] \leftarrow \{\text{random}(p), t\}$ 
7: COMMUNICATE( $\text{broadcast}(S_i[i])$ )
8:  $V_i \leftarrow \text{COMMUNICATE}(\text{collect}())$ 
9: if  $S_i[i][0] = 1$  and  $\exists k : p_k$  is active then
10:   return 0
11: else
12:   return 1
13: end if
14: return 1

```

---

## 2.3 Analysis

Now we just need to show that the expected number of processes which survive using the heterogeneous poison pill algorithm is  $\Theta(\log n)$ .

**Claim 3.** Let  $S$  be the set of processes which flipped 0 and survived and let  $U$  be the union of all the elements that they found not waiting. Then for  $i \in U$ , every  $j \in t_i$  is in  $U$ .

*Proof.* Intuitively this should make sense as  $U$  is pretty comprehensive (we are taking the union of a bunch of stuff). Since a process  $p_i$  which flipped zero can only survive if it sees only other processes which lost, every  $p \in U$  obtained 0. There is some entry of the view  $V_i$  corresponding to  $p$  which has 0. Further it also has the list associated with  $p$ . This list is taken into account during the union thus  $p_i$  see all the elements in the list  $t$  of  $p$ .  $\square$

**Claim 4.** *If processor  $q$  completed the first broadcast call no later than  $p$  completed its first propagate call, then  $q$  will be included in the  $l$  list of  $p$ .*

**Claim 5.** *The maximum expected number of processors that flip 0 and survive is  $O(\log n) + O(1)$ .*

*Proof.* Let  $S$  be the set of  $z$  processes. And  $U$  be defined as above. By Claim 4 and Claim 3, it must be the case that if  $p \in U$  and  $q$  finished the first broadcast no later than  $p$ , then  $q \in U$ . Notice that if we order the processes by the time they finish the first broadcast, all processes in  $U$  must come before the processes not in  $U$ .

Since  $U$  forms a closed set and all processes in  $U$  flipped 0, the probability that each process flipped 0 is at most  $\left(1 - \frac{\log |U|}{|U|}\right)$ . The probability for all processes in  $U$  to flip 0 is  $\left(1 - \frac{\log |U|}{|U|}\right)^{|U|} = O\left(\frac{1}{|U|}\right)$ . Since  $z \in U$ , This is  $O\left(\frac{1}{z}\right)$ .  $\square$

**Claim 6.** *The maximum expected number of processors that flip 0 and survive is  $O(\log^2 n) + O(1)$ .*

*Proof.* Order the processes according to the time they finish their first broadcast. By property 4, the process ordered first has  $|t| \geq 1$ , the process ordered second has  $|t| \geq 2$  and so on. Since the probability of flipping 1 decreases as  $|t|$  increases, the best expectation achievable is

$$1 + \sum_{l=2}^n \frac{\log |t|}{|t|} \in O(\log^2 n).$$

$\square$

## Reference

Alistarh, D., Gelashvili, R., and Vladu, A. (2015). How to elect a leader faster than a tournament. In *Proceedings of the 2015 ACM Symposium on Principles of Distributed Computing*, pages 365–374. ACM.