

Assume the graph  $G = (V, E)$  isn't  $(f + 1)$ -connected. Then there exists a set of  $f$  processes which if removed would disconnect the graph. Let  $G'$  be the graph after these processes are removed. Let  $p_i$  be a process in  $G'$ , and let  $V_1$  (respectively,  $V_2$ ) be the set of processes in  $G'$  which do (respectively, do not) have a path to  $p_i$  in  $G'$ .

If there exists a path between a process  $p_k \in V_2$  and a process  $p_j \in V_1$ , then the concatenation of that path with the path from  $p_j$  to  $p_i$  would be a path from  $p_k$  to  $p_i$ , which is a contradiction because  $p_k \in V_2$ . Therefore no messages can pass between  $V_1$  and  $V_2$ .

Assume there exists a consensus algorithm that works in this case. If all processes in  $V_1$  have input 0, then the algorithm must require them to output 0, because of the possibility (which is impossible for  $V_1$  to refute) that all processes in  $V_2$  also have input 0. Similarly, all processes in  $V_2$  would have to output 1 if all of their inputs are 1. But then the algorithm would fail if all processes in  $V_1$  have input 0 and all processes in  $V_2$  have input 1, so no valid algorithm exists.

Now assume the graph  $G = (V, E)$  is  $(f + 1)$ -connected. For each process  $p_i$  with input  $x_i$ , let  $m_i$  be a message saying that  $p_i$ 's input value is  $x_i$ .

---

**Algorithm 1**  $p_i$ 's Instructions in Round  $r$

---

- 1: **if**  $r=1$  **then** broadcast  $m_i$ .
  - 2: **else** broadcast all messages  $m_j$  that you just received for the first time.
  - 3: **end if**
  - 4:  $n \leftarrow$  the number of processes  $p_j$  for which you've received  $m_j$ .
  - 5: **if**  $r > n + 10$  **then** output the minimum value of  $x_j$  that you've seen.
  - 6: **end if**      $\triangleright$  The +10 is insurance against off by one errors in my analysis.
- 

First I assume that each process executes the above instructions even after outputting a value, and then I prove that this assumption isn't necessary.

**Lemma 1.** *Let  $p_s$  and  $p_t$  be processes that don't crash. If a message  $m_i$  reaches  $p_s$ , then  $m_i$  will also reach  $p_t$ .*

*Proof.* Since  $G$  is  $(f + 1)$ -connected and there are at most  $f$  crashes, there exists a path  $P$  from  $p_s$  to  $p_t$  in which no processes crash. If  $m_i$  doesn't propagate along  $P$  after reaching  $p_s$ , it's because  $m_i$  found a faster way to reach  $p_t$ , and no process sends the same message twice.  $\square$

**Lemma 2.** *Let  $p_t$  be a process that outputs a value.  $p_t$  will receive no messages after doing so.*

*Proof.* Let  $S$  be the set of  $n$  processes that  $p_t$  knows of when it outputs a value. Assume there exists  $p_s \in V - S$  such that there exists a path  $P$  from  $p_s$  to  $p_t$  along which  $m_s$  can propagate (i.e. no process in  $P$  crashes before it can send  $m_s$  at the appropriate time.) Without loss of generality assume that  $P$  is non-self-intersecting, and therefore finite. Let  $p_\mu$  be the last process in  $P$  in  $V - S$ , and let  $P'$  be the subpath of  $P$  from  $p_\mu$  to  $p_t$ . All but one element of  $P'$  is in  $S$ , so  $|P'| \leq 1 + |S| = 1 + n$ . Since  $P'$  is traversable when  $m_s$  reaches

$m_\mu$ ,  $P'$  is also traversable at the beginning of the algorithm (except perhaps in some cases where there's a shorter path). Therefore  $m_\mu$  will reach  $p_t$  by Round  $n + 2$  or so, which is a contradiction, because  $p_t$  outputs a value around Round  $r + 10$  and  $p_\mu \notin S$ .  $\square$

**Lemma 3.** *The algorithm performs the same, regardless of whether a processes continues to run its pseudocode after outputting a value.*

*Proof.* This follows from Lemma 2, the fact that no process sends a message several rounds after receiving it, and the fact that a process which terminates won't have received any messages for the last few rounds (because of the +10).  $\square$

**Theorem 1.** *The algorithm terminates.*

*Proof.* Each process waits a finite number of rounds before outputting a value or learning about a new process, and there are finitely many processes.  $\square$

**Theorem 2.** *All processes that don't crash output the same value.*

*Proof.* Let  $p_i$  be a process that doesn't crash, and let  $S$  be the set of processes that  $p_i$  knows about when outputting a value.  $S$  is nonempty because  $p_i \in S$ . By Lemmas 1-3, the value of  $S$  doesn't depend on the choice of  $p_i$ . All processes output  $\min_{p_j \in S} x_j$ .  $\square$

**Theorem 3.** *If all processes have the same input value, then no other value is output.*

*Proof.* Trivial.  $\square$