

CSC 2221 HW3

Eric Di Tomasso

Model: Synchronous message passing, n processes p_0, \dots, p_{n-1} . Let P be the set of all processes. Each process p_i receives an input x_i .

1 Impossibility Result

Lemma 1. *It is impossible to solve Byzantine Agreement deterministically in a synchronous message passing network when up to f processes can be Byzantine, if the network is not $(2f + 1)$ -connected.¹*

Proof. Suppose there is such an algorithm. Since the network is not $(2f + 1)$ -connected, there are 2 sets of processes S_0 and S_1 such that there are at most $2f$ process disjoint paths between them in the network. Thus, there is a set of processes C of size at most $2f$ such that all messages sent between S_0 and S_1 pass through at least 1 process in C . Let C_0 and C_1 be two disjoint sets such that $C_0 \cup C_1 = C$, $|C_0| \leq f$, and $|C_1| \leq f$. Let β be an execution of the algorithm on the ring pictured in Figure 1. Note that in the ring, each node corresponds to a set of processes, and an edge between 2 nodes corresponds to the set of all edges joining processes between the two sets. The number within each node corresponds to the input value for each process in the set of processes that the node represents. Consider the following 3 scenarios.

Scenario 1: Let α_1 be an execution of the algorithm in a network in which all processes start with input 0, and all processes in C_1 are faulty. Assume processes in the set C_1 are sending to S_0 the messages sent in β by C_1 , and are sending to C_0 and S_1 the messages sent in β by C'_1 . This scenario is illustrated in Figure 2.

Scenario 2: Let α_2 be an execution of the algorithm in a network in which all processes in S_1 start with input 0, all other processes start with input 1, and all processes in C_0 are faulty. Assume C_0 is sending to processes in S_0 the messages sent in β by C'_0 , and is sending to all processes in C_1 and S_1 the same messages as those sent in β by C_0 . This scenario is illustrated in Figure 3.

Scenario 3: Let α_3 be an execution of the algorithm in a network in which all processes start with input 1, and all processes in the set C_1 are faulty. Assume processes in the set C_1 are sending to S_0 the messages sent in β by C'_1 , and are sending to all neighbours in C_0 and S_1 the messages sent in β by C_1 . This scenario is illustrated in Figure 4.

I now argue that $\alpha_1 \stackrel{S_1}{\sim} \alpha_2$. Notice that in α_1 , processes in S_0 have the same initial configuration and receive the same messages from C_1 as they do in β . The same claim holds for the sets C_0 and S_1 . Thus, for all correct processes, the execution α_1 is indistinguishable from the execution β . In α_2 , processes in S_0 have the same initial configuration and receive the same messages from C_0 as processes in S'_0 do in β , and processes in C_1 have the same initial configuration and receive the same messages from C_0 as C'_1 in β . Thus in α_2 , S_1 has the same initial configuration it does in β , and receives the same messages it does in β . This is due to the fact that in both executions, the actions of the set of Byzantine processes are defined with respect to β . Hence, $\alpha_1 \stackrel{S_1}{\sim} \beta \stackrel{S_1}{\sim} \alpha_2$, so we conclude that $\alpha_1 \stackrel{S_1}{\sim} \alpha_2$. A similar argument can be used to show $\alpha_2 \stackrel{S_0}{\sim} \alpha_3$. By validity, all non-faulty processes in Scenario 1 must decide 0. Because $\alpha_1 \stackrel{S_1}{\sim} \alpha_2$, all processes in S_1 must also decide 0 in α_2 , and by agreement all processes in S_0 must decide 0 in α_2 . Applying indistinguishability again, it can be observed that all processes in S_0 must therefore decide 0 in α_3 . This violates validity, and is a contradiction. □

¹The high level approach for this proof can be found in course notes at the following URL: <http://www.cs.yale.edu/homes/aspnes/pinewiki/ByzantineAgreement.html>




Figure 1: The ring of processes.




Figure 2: Scenario 1.




Figure 3: Scenario 2.




Figure 4: Scenario 3.

2 An Algorithm for f -Resilient Consensus in a $(2f + 1)$ -Connected Network

Assume each process knows the topology of the network. A modified version of the exponential algorithm from the textbook is used (section 5.2.4). Before running the exponential algorithm, each process deterministically computes the same set of $2f + 1$ node-disjoint paths between every pair of processes in the network. Next, information gathering begins.

Instead of gathering information for the first $f + 1$ rounds, information is gathered for $f + 1$ phases. A phase consists of n rounds. In the first round of phase 1, each correct process p_i begins sending its input value to all other processes. This is performed as follows. For each process $p_j \neq p_i$, p_i sends a set containing a triple (i, j, x_i) to the first node on the path of each of the $2f + 1$ node-disjoint paths connecting p_i and p_j that were determined prior to running the algorithm. If multiple messages must be sent along the same edge, they are sent together in a set. For the next $n - 1$ rounds of phase 1, each process p_i receives all messages on its inbuffers, and forwards each message (j, k, x_j) to the next node on the node-disjoint path connecting p_j and p_k that p_i is a part of. If p_i receives a message from a process whose source and target processes don't match with the node-disjoint path that p_i is a part of, p_i does not forward the message.

At the end of round n of phase 1, each correct process p_i counts the maximum number of times d that it received the same input value x from another process p_j . If $d \geq f + 1$, p_i associates the input x with process p_j . Otherwise, p_i associates a default value \perp with p_j . The information each process knows is stored in a tree in the same fashion as the original algorithm. In general, in the first round of phase r , each correct process sends the r^{th} level of its tree to all other processes, using the same forwarding method that was used in phase 1.

By Lemma 2, by the end of each phase each correct process successfully identifies the message sent to it by every other correct process at the beginning of that phase. After round n of phase $f + 1$, each correct process uses the same recursive majority vote that was used in the exponential algorithm to decide its output. Thus, the modified algorithm behaves in the same way as the exponential algorithm on a complete graph. Correctness follows from the proof of the exponential algorithm.

Lemma 2. *In each round, each correct process p_i correctly identifies the message sent by all other correct processes p_j to p_i .*

Proof. Each correct process p_j that sends some information v to p_i sends it along $2f + 1$ node-disjoint paths. Since there are at most f Byzantine processes, there are at least $f + 1$ non-faulty paths. Since the length of each non-faulty path is at most $n - 1$, by the beginning of round n the messages on all non-faulty node-disjoint paths from p_j to p_i will have arrived to p_i , and p_i will have received at least $f + 1$ copies of v . Thus p_i will identify v as the message sent by p_j that round. \square